

Part 4

Thread Synchronization Primitives

Synchronization Options

- The threading library defines the following objects for synchronizing threads
 - Lock
 - RLock
 - Semaphore
 - BoundedSemaphore
 - Event
 - Condition

Synchronization Options

- In my experience, there is often a lot of confusion concerning the intended use of the various synchronization objects
- Maybe because this is where most students "space out" in their operating system course (well, yes actually)
- Anyways, let's take a little tour

Mutex Locks

- Mutual Exclusion Lock

`m = threading.Lock()`

- Probably the most commonly used synchronization primitive
- Primarily used to synchronize threads so that only one thread can make modifications to shared data at any given time

Mutex Locks

- There are two basic operations

```
m.acquire()  
m.release()
```

```
# Acquire the lock  
# Release the lock
```

- Only one thread can successfully acquire the lock at any given time
- If another thread tries to acquire the lock when its already in use, it gets blocked until the lock is released

Use of Mutex Locks

- Commonly used to enclose critical sections

```
x = 0
x_lock = threading.Lock()
```

Thread-1

```
...
x_lock.acquire()
```

Critical
Section

```
x = x + 1
```

```
x_lock.release()
```

```
...
```

Thread-2

```
...
x_lock.acquire()
```

```
x = x - 1
```

```
x_lock.release()
```

```
...
```

- Only one thread can execute in critical section at a time (lock gives exclusive access)

Using a Mutex Lock

- It is your responsibility to identify and lock all "critical sections"

```
x = 0
x_lock = threading.Lock()
```

Thread-1

```
...
x_lock.acquire()
x = x + 1
x_lock.release()
...
```

Thread-2

```
...
x = x - 1
...
```



If you use a lock in one place, but not another, then you're missing the whole point. All modifications to shared state must be enclosed by lock acquire()/release().

Locking Perils

- Locking looks straightforward
- Until you start adding it to your code
- Managing locks is a lot harder than it looks

Lock Management

- Acquired locks must always be released
- However, it gets evil with exceptions and other non-linear forms of control-flow
- Always try to follow this prototype:

```
x = 0
x_lock = threading.Lock()

# Example critical section
x_lock.acquire()
try:
    statements using x
finally:
    x_lock.release()
```

Lock Management

- Python 2.6/3.0 has an improved mechanism for dealing with locks and critical sections

```
x = 0
x_lock = threading.Lock()

# Critical section
with x_lock:
    statements using x
...
```

- This automatically acquires the lock and releases it when control enters/exits the associated block of statements

Locks and Deadlock

- Don't write code that acquires more than one mutex lock at a time

```
x = 0
y = 0
x_lock = threading.Lock()
y_lock = threading.Lock()
```

```
with x_lock:
    statements using x
...
with y_lock:
    statements using y
...
```

- This almost invariably ends up creating a program that mysteriously deadlocks (even more fun to debug than a race condition)

RLock

- Reentrant Mutex Lock

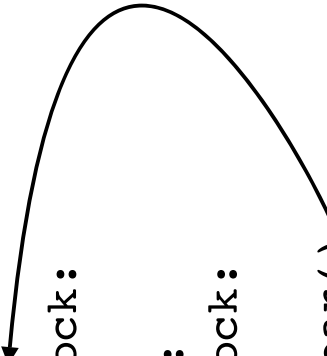
```
m = threading.RLock()      # Create a lock
m.acquire()                 # Acquire the lock
m.release()                 # Release the lock
```

- Similar to a normal lock except that it can be reacquired multiple times by the same thread
- However, each `acquire()` must have a `release()`
- Common use : Code-based locking (where you're locking function/method execution as opposed to data access)

RLock Example

- Implementing a kind of "monitor" object

```
class Foo(object):  
    lock = threading.RLock()  
    def bar(self):  
        with Foo.lock:  
            ...  
    def spam(self):  
        with Foo.lock:  
            ...  
            self.bar()  
            ...
```



- Only one thread is allowed to execute methods in the class at any given time
- However, methods can call other methods that are holding the lock (in the same thread)

Semaphores

- A counter-based synchronization primitive

```
m = threading.Semaphore(n) # Create a semaphore
m.acquire()                # Acquire
m.release()                # Release
```
- `acquire()` - Waits if the count is 0, otherwise decrements the count and continues
- `release()` - Increments the count and signals waiting threads (if any)
- Unlike locks, `acquire()/release()` can be called in any order and by any thread

Semaphore Uses

- Resource control. You can limit the number of threads performing certain operations. For example, performing database queries, making network connections, etc.
- Signaling. Semaphores can be used to send "signals" between threads. For example, having one thread wake up another thread.

Resource Control

- Using a semaphore to limit resources

```
sema = threading.Semaphore(5)    # Max: 5-threads


def fetch_page(url):
    sema.acquire()
    try:
        u = urllib.urlopen(url)
        return u.read()
    finally:
        sema.release()
```

- In this example, only 5 threads can be executing the function at once (if there are more, they will have to wait)

Thread Signaling

- Using a semaphore to signal

`done = threading.Semaphore(0)`



- Here, `acquire()` and `release()` occur in different threads and in a different order
- Often used with producer-consumer problems

Events

- Event Objects

```
e = threading.Event()  
e.isSet()      # Return True if event set  
e.set()        # Set event  
e.clear()      # Clear event  
e.wait()       # Wait for event
```

- This can be used to have one or more threads wait for something to occur
- Setting an event will unblock all waiting threads simultaneously (if any)
- Common use : barriers, notification

Event Example

- Using an event to ensure proper initialization

```
init = threading.Event()

def worker():
    init.wait()      # Wait until initialized
    statements
    ...

def initialize():
    statements      # Setting up
    statements      # ...
    ...
    init.set()      # Done initializing

Thread(target=worker).start()      # Launch workers
Thread(target=worker).start()
Thread(target=worker).start()
initialize()                        # Initialize
```

Event Example

- Using an event to signal "completion"

```
def master():  
    ...  
    item = create_item()  
    evt = Event()  
    worker.send((item,evt))  
    ...  
    # Other processing  
    ...  
    ...  
    ...  
    ...  
    ...  
    # Wait for worker  
    evt.wait()
```

Worker Thread

```
item, evt = get_work()  
processing  
processing  
...  
...  
# Done  
evt.set()
```

- Might use for asynchronous processing, etc.

Condition Variables

- Condition Objects

```
cv = threading.Condition([lock])
cv.acquire()      # Acquire the underlying lock
cv.release()      # Release the underlying lock
cv.wait()         # Wait for condition
cv.notify()       # Signal that a condition holds
cv.notifyAll()    # Signal all threads waiting
```

- A combination of locking/signaling
- Lock is used to protect code that establishes some sort of "condition" (e.g., data available)
- Signal is used to notify other threads that a "condition" has changed state

Condition Variables

- Common Use : Producer/Consumer patterns

```
items = []  
items_cv = threading.Condition()
```

Producer Thread

```
item = produce_item()  
with items_cv:  
    items.append(item)
```

Consumer Thread

```
with items_cv:  
    ...  
    x = items.pop(0)  
  
# Do something with x  
...
```

- First, you use the locking part of a CV
synchronize access to shared data (items)

Condition Variables

- Common Use : Producer/Consumer patterns

```
items = []  
items_cv = threading.Condition()
```

Producer Thread

```
item = produce_item()  
with items_cv:  
    items.append(item)  
    items_cv.notify()
```

Consumer Thread

```
with items_cv:  
    while not items:  
        items_cv.wait()  
    x = items.pop(0)  
  
    # Do something with x  
    ...
```

A curved arrow originates from the `items_cv.notify()` line in the Producer Thread and points to the `items_cv.wait()` line in the Consumer Thread, illustrating the signaling mechanism.

- Next you add signaling and waiting
- Here, the producer signals the consumer that it put data into the shared list

Condition Variables

- Some tricky bits involving `wait()`
- Before waiting, you have to acquire the lock
- `wait()` releases the lock when waiting and reacquires when woken
- Conditions are often transient and may not hold by the time `wait()` returns. So, you must always double-check (hence, the while loop)

Consumer Thread

→ with `items_cv:`

while not items:

→ `items_cv.wait()`

`x = items.pop(0)`

Do something with x

...

Interlude

- Working with all of the synchronization primitives is a lot trickier than it looks
- There are a lot of nasty corner cases and horrible things that can go wrong
- Bad performance, deadlock, livelock, starvation, bizarre CPU scheduling, etc...
- All are valid reasons to not use threads